

AI Boot Camp **Final Project**

Retrieval-Augmented Generation (RAG) LLM

Team Members:

Jeff Flachman

Kerek Spinney

Ashwini Kumar

Exec Summary

Purpose:

Local Large Language Model (LLM) capable of reviewing, summarizing, and leveraging proprietary documents to facilitate a local knowledge base.

Key Features:

- Utilize RAG (Retrieval-Augmented Generation) Workflow to create and query the datastore of documents
 - Loads documents (e.g. pdf, word, etc) into a vector database
 - Use LLM answer query based on the content from the data files
 - Provides a summary of the content based on a query
- Quick knowledge base retrieval
- Can be deployed on a laptop or url
- Scalable

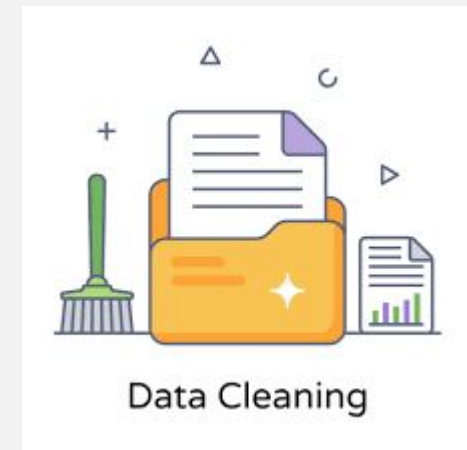
Goals

1. Build Prototype version using RAG and LLM
2. Load data related to “GenAI”
3. Give the user the ability to ask a knowledge base questions and return with accurate answers

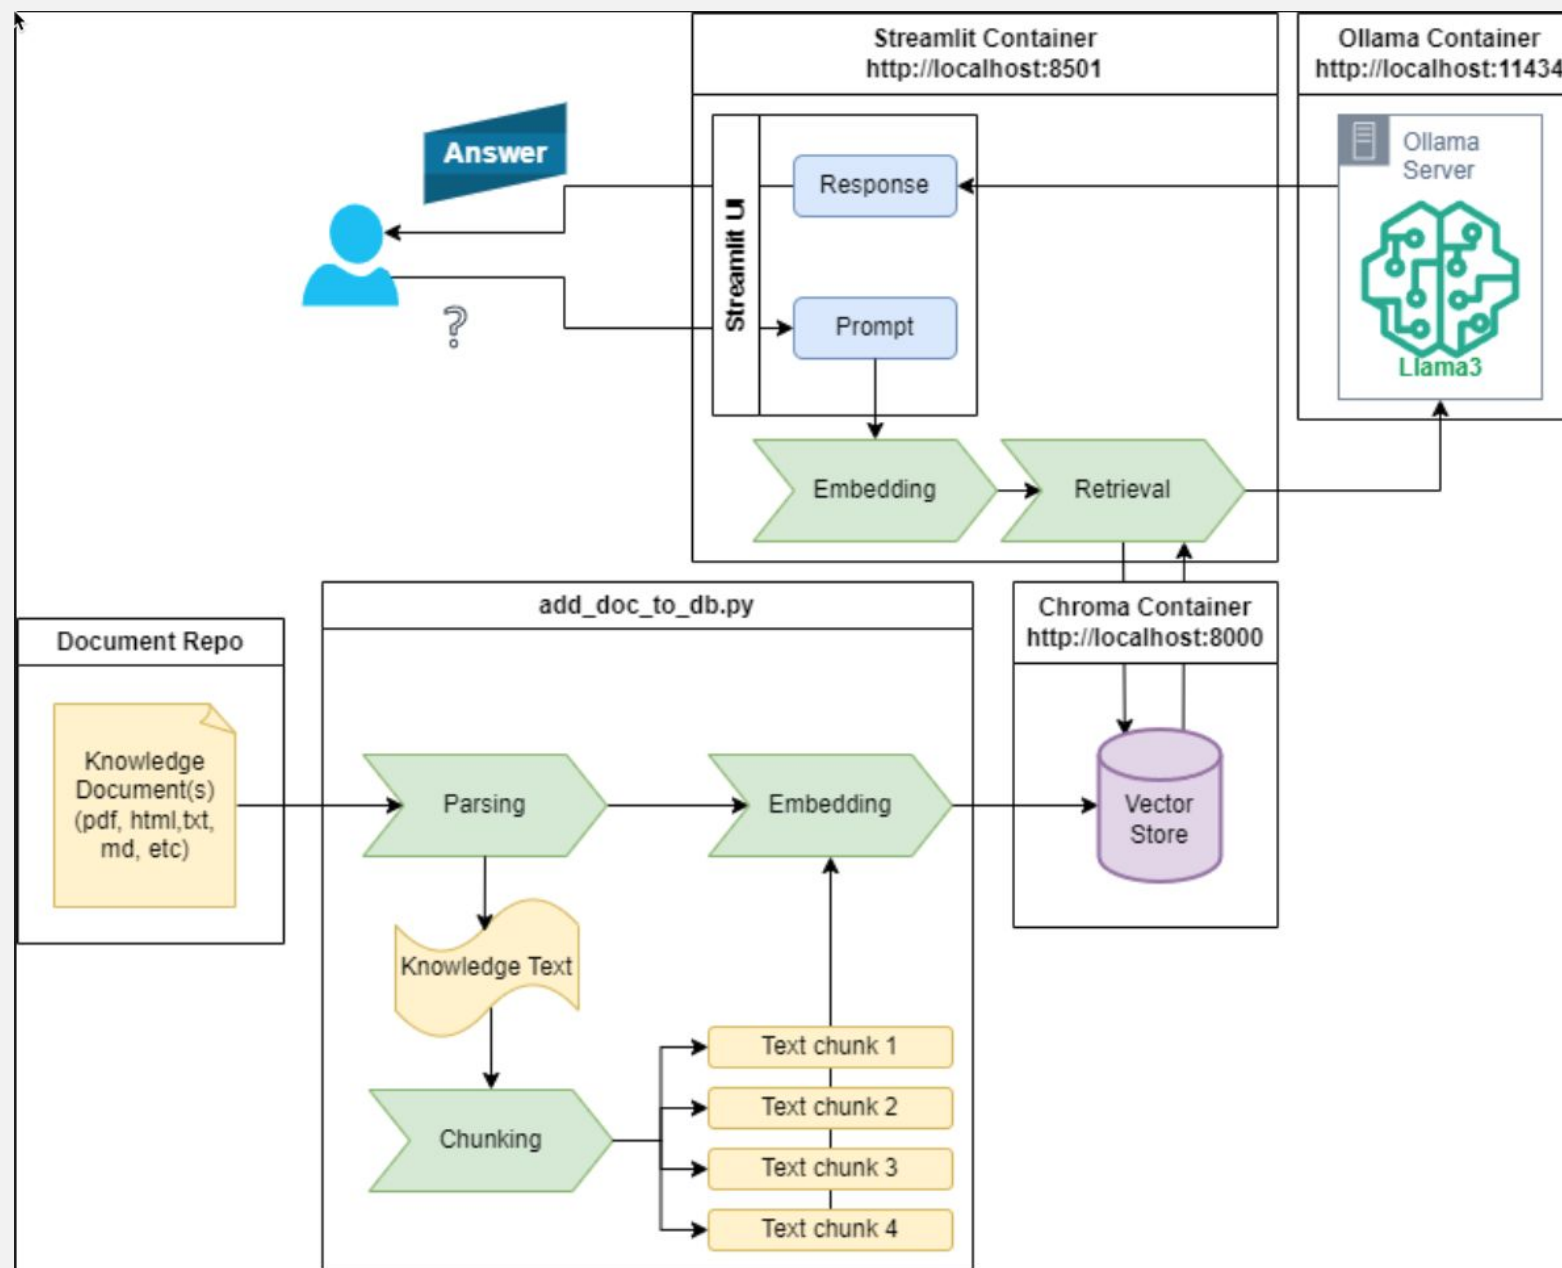


Data Collection & Cleaning

1. **Researched and Gathered** information related to GenAI.
2. **Saved all content as PDF** and loaded into Data folder
3. **PDFs are a graphic method for printing a document.**
 - a. They do not contain attributes like header, footer, footnotes, etc.
 - b. These artifacts end up in the middle of paragraph the spans two pages.
 - c. This greatly impacts the effectiveness of a RAG solution.
4. **Text files and Word Documents are better formats** to import with fewer document artifacts impacting the effectiveness of the solution



Technical Approach



Model Optimization

Importing documents:

- **Chunk/Splitting the text can be done in many different ways.**
 - **Word Tokens:** Or random groups of words is too Simple
 - **Implemented:** Split into sentences as chunks
 - **Future Goal:** Semantic splitting groups sentences that talk about the same type of information using embeddings of each sentence group

LLM Selection:

- **Trade off between model size and model robustness**
- **Used Huggingface Leaderboard:** to select the best performing model for a model size.
- **Running on a local laptop limits the size of the models that can be used.**
 - More performance models are ~0.5 Billion parameters model vs 7b or 70b.
 - Selected 0.5b model to provide perfect response on local system.

Demo

DEMO

Challenges Encountered

- **Setting up Docker.**
- **Finding materials online to assist in our project.**
 - Sort through many different approaches that did not necessarily meet our deployment methods.
- **Combining all portions of the project and have it run.**
- **Taking on a large project.**
- **Using PDFs.**



Future Considerations



- Test other LLM model options.
- Test other chunking/splitting functions.
- Test additional data and file types
- Work with Streamlit to add a user interface.
- Secure access





Questions?

Technical Approach

- **Python file** to read, convert, chunk and load documents into the Vector Database
- **Vector Database:** used ChromaDB in a docker container to provide easy configuration, deployment and use of the vector database
- **User Interface python file**
 - **Uses Streamlet** to allow users to prompt questions that are answered based on the documents loaded into the vector database
 - **Uses the RAG workflow**
 - Take a prompt, and get closest matching document chunks from ChromaDB
 - Submit the Chunks as context along with a system context and the users prompt to the LLM
 - Displays the Response based on the document chunks back to the user.
- **LLM Server** using llama-cpp-server in a docker container that loads any number of LLM models for use in the solution.